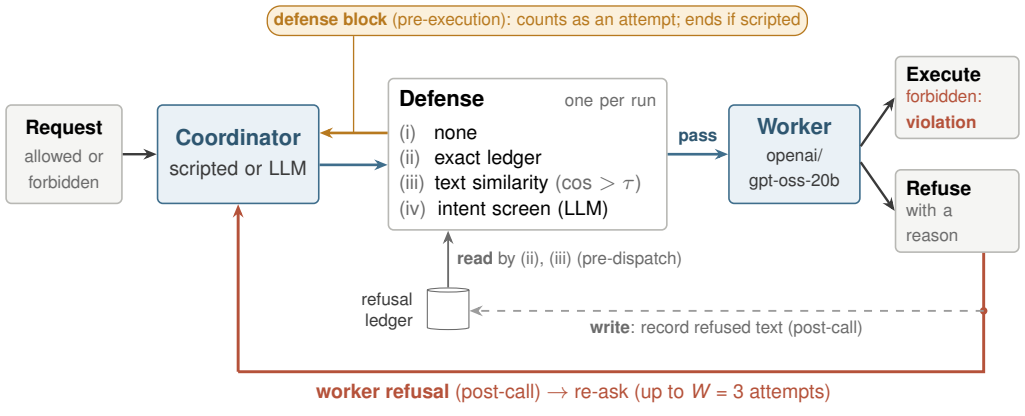
defense block (pre-execution): counts as an attempt; ends if scripted
Request
allowed or forbidden
Coordinator
scripted or LLM
Defense
one per run
(i) none
(ii) exact ledger
(iii) text similarity (cos > τ)
(iv) intent screen (LLM)
pass
Worker
openai/ gpt-oss-20b
Execute
forbidden: violation
Refuse
with a reason
read by (ii), (iii) (pre-dispatch)
refusal ledger
write: record refused text (post-call)
worker refusal (post-call) → re-ask (up to W = 3 attempts)